

# Device Drivers and Memory Management

## Linux Device Drivers

### User side

Device drivers have a special file in `/dev` directory and have 5 system calls

- `Open`: Make device available
- `read`: Read from device
- `write`: Write to device
- `ioctl`: Input Output Control (Optional)
- `close`: Make device unavailable

### Kernel side

Each file has functions associated with it and are called when corresponding system calls made

`linux/fs.h` shows all ops on files

Device drivers implement at least `open`, `read`, `write`, `close`

## Categorisation

Kernel tracks:

**Physical Dependencies** between devices (E.g. devices on USB hub)

**Buses** Channels between processor and devices. Physical or logical

**Classes** Sets of same type devices

## Handling interrupts

Cycle for interrupt handling:

- Device sends interrupt
- CPU selects interrupt handler
- Handler processes interrupt to do two tasks:
  - Data transferred between device
  - Wake up processes waiting for transfer to finish
- Handler clears interrupt bit of device for next interrupt

Interrupt processing must be as short as possible

Data transfer fast but processing slow

Interrupt processing has 2 halves:

**Top half** called directly by handler and transfers data between device and kernel buffer and schedules software interrupt to start bottom half

**Bottom half** runs in interrupt context and does rest of processing (eg working thru protocol stack, wake up processes)

## Memory Management

Memory is a **limited resource** and memory hunger of applications increases with more memory

Requires **sophisticated algorithms**, with support from HW, compiler and loader

Both programs (**logical**) and HW (**physical**) view memory the same, a set of memory cells starting at 0x0 and ending at some value. Problem occurs when we want to store multiple programs simultaneously, all who views memory as starting at 0x0. This requires a *mapping between logical and physical addresses*

This mapping can occur at different times:

**Compile time**

*absolute references* generated

**Load time**

mapping done by *special program (loader)*

**Execution time**

*HW support*

Address mapping can take a step further with **dynamic linking**, allowing a *single copy of a system library to be shared among other programs*. Requires OS support for code to be accessible to multiple processes

## Swapping

If **memory demand too high**, memory for some processes **transferred to disk**

With *scheduling*, low priority processes are swapped out

Problems:

**Long transfer times**

**Pending IO?**

Therefore, *swapping is not a principle memory management technique*, meaning is not a solution for management

## Fragmentation

Swapping creates problems:

- Over time, **small holes** appear in memory (*external frag*)
- Programs smaller than those holes do not fill them and **leftover space too small to be a hole** (*internal frag*)

# Fragmentation



Assume hole must be at least 4 blocks to be considered a "hole"

Now a new program is put in but only actually takes up 3 holes



Too small to be considered a "hole"

**Strategies** for choosing holes:

*First fit* Start from beginning and use first available hole

*Rotating first fit* start after last assigned part of memory

*Best fit* find smallest usable hole

*Buddy system* Free holes administered according to tree structure; smallest chunk used

## Paging

Approach:

Assign **memory of fixed size** (*page*) which avoids *external frag*

*Page table translates* logical to physical addresses

**HW support mandatory:**

If page table small, use *fast registers*. Store large tables in main mem but *cache entries*. This is the **general principle**

**Memory protection** easily added as protection info stored in table

# Segmentation

Divide memory according to usage:

- Data* mutable, different foreach instance
- Program Code* immutable, same foreach instance
- Symbol Table* immutable, same foreach instance (*debugging*)

Requires HW support:

Same as *paging* but overflow check needed

*Paging* for ease of allocation, *segmentation* for memory use so most systems use both

## Virtual memory

Completely separates logical and physical memory and programs to use a lot of memory

Works as most programs don't use that much

Efficient implementation tricky as difference between memory and disk access immense

## Demand Paging

Virtual memory implemented as *demand paging*: memory divided into *pages of same length*, with *validation bit*

Decisions to make:

- Which processes to *swap out* which moves memory to disk and blocks process (swapper)
- Which *pages to move to disk* when new page required (pager)

Page fault rate must be minimised (page has to be fetched from disk) **Page Fault handling**

## Page replacement algos

FIFO



- easy to implement



- locality not considered
- Belady's anomaly: Increase in number of frames increases number of page faults

NOT EXPLAINED BUT PAGE FRAMES = PHYSICAL MEMORY, PAGE = VIRTUAL MEMORY

Optimal algorithm

Replace pages which won't be used or used last



- Nice to compare



- Unimplementable

Least recently used

Replace page which has been unused the longest



-Likely to have low page fult rate



-Requires a lot of HW support

Possible HW solutions:

*Use a stack in "microcode"* (stack within the MMU or HW itself)

*Reference bit approximation*: HW sets bits to 1 when page referenced, FIFO through pages and use first page with bit=0 and reset pages with bit=1 to 0 as you go along *second chance algo*

## Thrashing

If a process lacks **frames it uses constantly**, page fault rate high  
*CPU throughput decreases thus negatively affects performance*

Solutions:

*Working set model* (locality based)

A *Locality* is a set of pages used together, locality model assumes programs execute move from locality to locality.

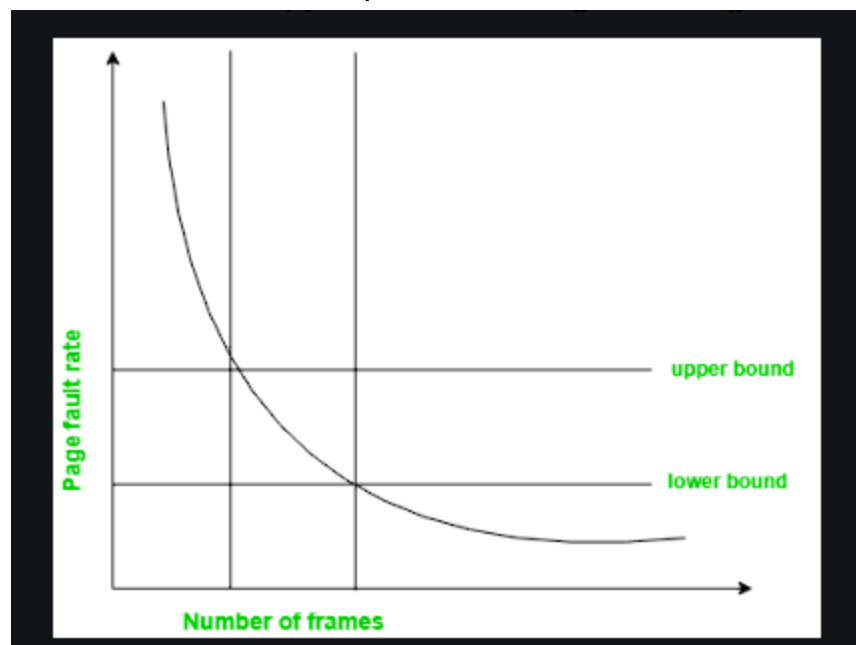
Define *working set* as pages used in most recent "X" page references and keep these in memory

Achieves high CPU utilisation and prevents thrashing if  $X \geq$  current locality size

*Page-Fault Frequency*

-Give process more frames if page frequency rate high

-remove frame from process if fault rate low



## Memory management in linux kernel

4 Segments:

*Kernel code*

*Kernel data*

*User code*

*User data*

Paging used with permission system

# Kernel and user memory

Separate *logical* addresses for kernel and user memory

For 32bit arch

kernel space is upper 1 GB

user space is lower 3 GB

kernel memory always mapped but protected from user process access

For 64 bit arch

kernel space is upper half

user space is lower half

## Page caches (Prolly won't be tested)

Often, there are repeated cycles of alloc and free of similar objects (inodes, dentries)

Pool of pages used as a cache *slab cache* maintained by application (eg file system)

*kmalloc* uses slab cache for commonly used sizes

## Summary

### Device Drivers

implement *open*, *read*, *write*, *close*, with common structure

### Mem management

- Serious effort to manage limited resource
- Isolate mem for each process
- If mem demand high, swap out parts of process memory
- Paging* and *segmentation*
- Requires HW support