

# File systems

# File System

- Function: main permanent data storage  
**Speed bottleneck!**
- **Capacity not a problem** nowadays: 2 TB disks even for PC.  
But **backup becoming a problem**.
- Logical view (view of programmer): **tree structure of files** together with read/write operation and creation of directories
- Physical view: **sequence of blocks**, which can be read and written. OS has to map logical view to physical view, must impose tree structure and assign blocks for each file

Two main possibilities to realize filesystem:

- **Linked list**: Each block contains pointer to next  
⇒ Problem: random access (`seek()`) costly: have to go through whole file until desired position.
- **Indexed allocation**: Store pointers in one location: so-called index block (similar to page table). To cope with vastly differing file sizes, may introduce **indirect index blocks**.

Index blocks are called `inodes` in Unix.

Inodes store additional information about the file (eg size, permissions)

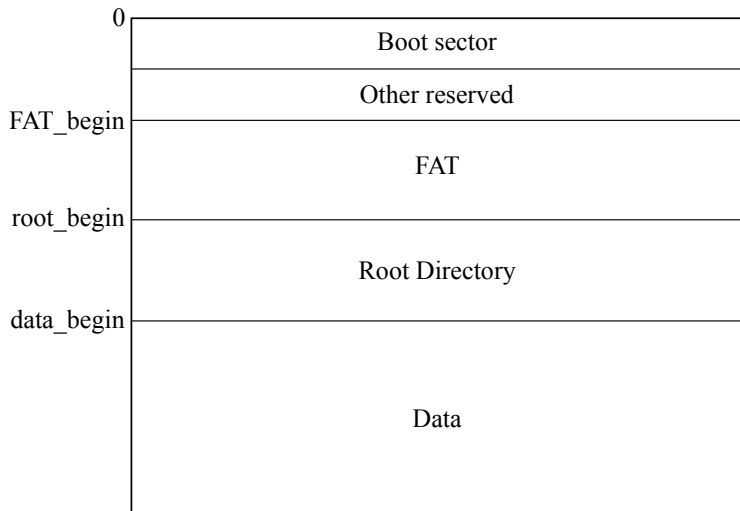
## Worked Example

Worked example – based on  
<http://www.tavi.co.uk/phobos/fat.html>

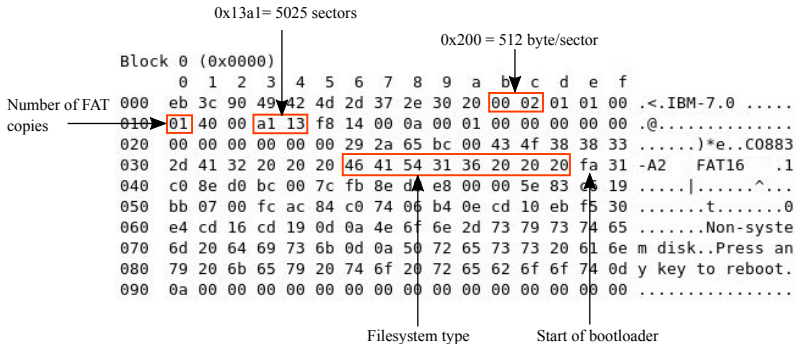
## Example: FAT

- F(ile) A(llocation) T(able) – dates back to 70s.
- Useful for explaining filesystem concepts, modern filesystems are more complicated
- Variants FAT12. FAT16, FAT32 define number of bits per FAT entry – we focus on FAT16
- Sector = disk unit (e.g. 512 byte), aka block
- Cluster = multiple sectors (factor 1, 2, 4, ..., 128)  
(*here*: assume cluster = 1 sector)
- Uses linked list (“cluster chain”) to group clusters

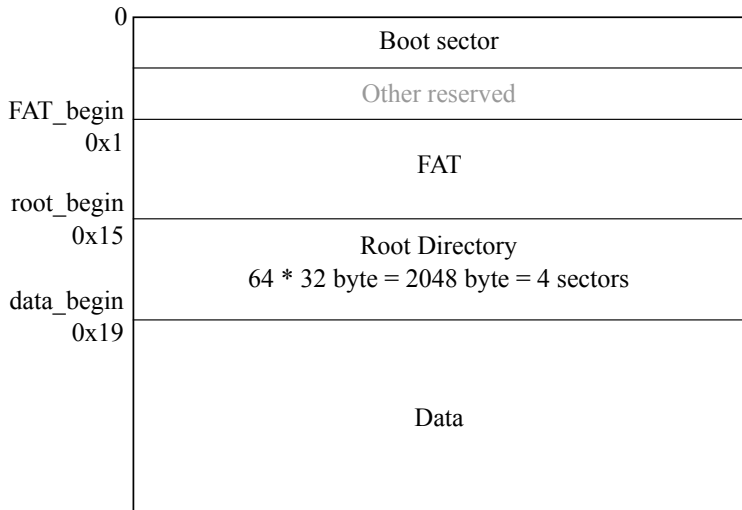
## Example: FAT16 Structure



# Example: FAT16 Bootsector



## Example: FAT16 with Offsets





# Example: File Allocation Table

FAT\_begin

FFF0	FFFF	0003 <sub>2</sub>	0004 <sub>3</sub>
0006 <sub>4</sub>	0000 <sub>5</sub>	FFFF <sub>6</sub>	0008 <sub>7</sub>
0009 <sub>8</sub>	FFFF <sub>9</sub>	0000 <sub>A</sub>	0000 <sub>B</sub>
000F <sub>C</sub>	0010 <sub>D</sub>	0000 <sub>E</sub>	000D <sub>F</sub>
FFFF <sub>10</sub>	0000 <sub>12</sub>	0000 <sub>13</sub>	0000 <sub>14</sub>

2 → 3 → 4 → 6      7 → 8 → 9      C → F → D → 10

# Example: File in Root Directory

Block 21 (0x0015)

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	
000	43	4f	38	38	33	2d	41	32	20	20	20	28	00	00	00	00	C0883-A2 (....
010	00	00	00	00	00	00	91	9e	65	39	00	00	00	00	00	00	.....e9.....
020	46	4f	4f	42	41	52	20	20	54	58	54	21	00	a3	91	9e	FOOBAR TXT!....
030	65	39	65	39	00	00	91	9e	65	39	c6	10	1a	00	00	00	e9e9....e9.....
040	4e	45	54	57	4f	52	4b	20	56	52	53	20	00	b6	91	9e	NETWORK VRS ....
050	65	39	65	39	00	00	91	9e	65	39	4e 0f	92 06 00 00					e9e9...e9N.....
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....

sector 0xF4E

length 0x692 = 1682 byte

filename & extension

# Example: File in FAT

Block 16 (0x0010)

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
000	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
010	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
090	00	00	00	00	00	00	00	00	00	00	00	00	4f	0f	50	0f
0a0	51	0f	ff	ff	00	00	00	00	00	00	00	00	00	00	00	00
0b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

4 block cluster  
chain (2048 byte)

## Example: FAT Limits

- Max. volume size: 2 GB ( $2^{16} \cdot 32 \text{ kB}$ )
- Max. file size: 2 GB
- Max. number of files: 65,460 (32 kB clusters)
- Max. filename length: 8 + 3
- FAT32 / exFAT have higher limits
- Newer filesystems (NTFS, ext4) also overcome these limits, using other data structures (e.g. B-tree for dir structure, bitmap for allocation)

## Further Aspects

### Further aspects of filesystems

# Caching

Disk blocks used for storing directories or recently used files cached in main memory

Blocks periodically written to disk

⇒ Big efficiency gain

Inconsistency arises when system crashes

Reason why computers must be shutdown properly

# Journaling File Systems

To minimise data loss at system crashes, ideas from databases are used:

- Define **Transaction points**: Points where cache is written to disk  
⇒ Have consistent state
- Keep log-file for each write-operation  
Log enough information to unravel any changes done after latest transaction point

# Disk Access

Disk access contains three parts:

- **Seek**: head moves to appropriate track
- **Latency**: correct block is under head
- **Transfer**: data transfer

HDDs: Time necessary for seek and latency dwarfs transfer time

⇒ Distribution of data and scheduling algorithms have vital impact on performance for HDDs, less so for SSDs



# Disk Scheduling Algorithms

Standard algorithms apply, adapted to the special situation:

1.) **FCFS**: easiest to implement, but: may require lots of head movements

2.) **Shortest Seek Time First**: Select job with minimal head movement

Problems:

- may cause starvation
- Tracks in the middle of disk preferred

Algorithm does not minimise number of head movements

3.) **SCAN-scheduling**: Head continuously scans the disk from end to end (lift strategy)

⇒ solves the fairness and starvation problem of SSTF

Improvement: **LOOK-scheduling**:

head only moved as far as last request (lift strategy).

Particular tasks may require different disk access algorithms

Example : Swap space management

Speed absolutely crucial ⇒ different treatment:

- Swap space stored on **separate partition**
- **Indirect access methods not used**
- **Special algorithms used for access of blocks**  
Optimised for speed at the cost of space (eg increased internal fragmentation)

# Linux Implementation

Interoperability with Windows and Mac requires support of different file systems (eg vfat)

⇒ Linux implements common interface for all filesystems

Common interface called **virtual file system**

virtual file system maintains

- inodes for files and directories
- caches, in particular for directories
- superblocks for file systems

All system calls (eg open, read, write and close) first go to virtual file system

If necessary, virtual file system selects appropriate operation from real file system

# Disk Scheduler

Kernel makes it possible to have different schedulers for different file systems

Default scheduler (Completely Fair Queuing) based on lift strategy  
have in addition separate queue for disk requests for each process  
queues served in Round-Robin fashion

Have in addition No-op scheduler: implements FIFO

Suitable for SSD's where access time for all sectors is equal